



What are the major limitations of Docker that Kubernetes overcomes?

Key Problems Docker Has, That Kubernetes Solves

1 Single Host Limitation

- **Docker** (when used alone, without orchestration tools) runs containers on a **single machine**. This means all containers are limited to one host.
- **Kubernetes** runs containers across **multiple machines (a cluster)**, allowing for better resource utilization, redundancy, and scalability.

2 Auto Scaling

- **Docker** requires **manual scaling**—you have to manually spin up more containers using commands like `docker run` or `docker-compose scale`.
- **Kubernetes** has **Horizontal Pod Autoscaling (HPA)**, which can automatically increase or decrease the number of running container instances based on CPU, memory, or custom metrics.

3 Auto Healing

- **Docker** does not automatically restart failed containers (unless you use something like `docker restart policies`).
- **Kubernetes** has a built-in **self-healing mechanism**:
 - ✓ It restarts failed containers
 - ✓ It replaces unresponsive containers
 - ✓ It reschedules containers if a node fails (Kubernetes **automatically reschedules** the containers (Pods) running on that node to **other available healthy nodes** in the cluster.)

4 Enterprise Support & Management

- **Docker (Standalone)** is not built for large-scale, enterprise-level deployments. It lacks built-in features for **load balancing, service discovery, rolling updates, or advanced security policies**.
- **Kubernetes** is designed with **enterprise needs** in mind:
 - ✓ Supports **RBAC (Role-Based Access Control)** for security
 - ✓ Offers **rolling updates & rollbacks**
 - ✓ Has **built-in networking & service discovery**
 - ✓ Works with **multi-cloud & hybrid cloud environments**

Summary

Docker is great for **containerization**, but **Kubernetes is a full-fledged orchestration system** that solves real-world problems in **scalability, high availability, and automation**.

Q1: What is a key limitation of using Docker without orchestration tools?

- A) It can only run on Linux-based systems
- B) Containers are limited to a single machine
- C) It does not support container networking
- D) Docker automatically scales containers across multiple machines

✓ **Correct Answer: B**

Q2: How does Kubernetes handle scaling differently from standalone Docker?

- A) Kubernetes provides automatic scaling based on resource usage
- B) Docker automatically scales containers without user intervention
- C) Kubernetes does not support auto-scaling
- D) Docker can scale across multiple hosts without any additional tools

✓ **Correct Answer: A**

Q3: What is one advantage of Kubernetes' self-healing mechanism over Docker?

- A) Kubernetes does not restart failed containers
- B) Docker automatically replaces unresponsive containers
- C) Kubernetes reschedules containers if a node fails
- D) Docker ensures all containers remain running without external tools

✓ **Correct Answer: C**

Q4: Which of the following is a reason why enterprises prefer Kubernetes over standalone Docker?

- A) Kubernetes has built-in support for RBAC and rolling updates
- B) Docker offers better load balancing for large-scale deployments
- C) Kubernetes does not support multi-node container management
- D) Standalone Docker provides better security features than Kubernetes

✓ **Correct Answer: A**