



How to Safely Drain a Kubernetes Node?

Draining a node means **safely evicting all Pods** from a node before maintenance, scaling down, or decommissioning. To ensure **zero downtime** and minimal disruption, follow these steps:

📌 Step 1: Cordon the Node

Before draining, **cordon** the node to prevent new Pods from being scheduled on it.

```
sh
```

```
kubectl cordon <node-name>
```

◆ What does this do?

- Marks the node as **unschedulable**
 - Existing Pods keep running, but **no new Pods** will be scheduled here
-

📌 Step 2: Drain the Node

Now, safely evict Pods using:

```
sh
```

```
kubectl drain <node-name> --ignore-daemonsets --delete-emptydir-data
```

◆ Breakdown of options:

- `--ignore-daemonsets`: Skips **DaemonSet Pods** (e.g., logging, monitoring Pods like `fluentd`, `kube-proxy`).
- `--delete-emptydir-data`: Ensures `emptyDir` volumes don't block eviction.

- `--force` (*Optional*): Use only if Pods without controllers block eviction.

⚠ Common Errors & Fixes:

- **Error: "Cannot delete Pods that declare no controller"**
 - ◆ Use `--force` (not recommended unless necessary).

```
sh
-----
kubectl drain <node-name> --ignore-daemonsets --delete-emptydir-data --force
```

- **Error: "PDB prevents eviction"**
 - ◆ Check the Pod Disruption Budget (PDB):

```
sh
-----
kubectl get pdb -A
```

If `MIN AVAILABLE` isn't met, Kubernetes won't drain Pods until a replacement is ready.

📌 Step 3: Verify the Node is Drained

Check that the node has **no more scheduled Pods** (except DaemonSets):

```
sh
-----
kubectl get pods --all-namespaces --field-selector spec.nodeName=<node-name>
```

If empty, the node is successfully drained.

📌 Step 4: Uncordon the Node (After Maintenance)

If you want to **bring the node back**, uncordon it so new Pods can be scheduled:

```
sh
-----
kubectl uncordon <node-name>
```

What is the main purpose of draining a node in Kubernetes?

- A) To permanently delete all Pods running on the node
- B) To safely evict all Pods so that the node can be **rebooted, upgraded, or removed**
- C) To delete the node from the cluster
- D) To restart the Kubernetes cluster

✅ **Correct Answer: B**

Command to Drain a Node

Which command is used to safely drain a node before maintenance?

- A) `kubectrl delete node <node-name>`
- B) `kubectrl drain <node-name> --ignore-daemonsets`
- C) `kubectrl stop node <node-name>`
- D) `kubectrl cordon <node-name>`

 **Correct Answer: B**

Preventing New Pods from Scheduling

Before draining a node, how can you **prevent new Pods** from being scheduled on it?

- A) `kubectrl drain <node-name>`
- B) `kubectrl delete node <node-name>`
- C) `kubectrl cordon <node-name>`
- D) `kubectrl stop <node-name>`

 **Correct Answer: C**