



## What are Resource Requests and Limits in Kubernetes?

In Kubernetes, each container in a Pod can declare:

- **Resource Requests** → *Minimum* amount of CPU/memory guaranteed to the container.
  - **Resource Limits** → *Maximum* amount of CPU/memory a container is allowed to use.
- 

 **Think of it like booking food at a buffet:**

- **Request** = You book at least 1 plate of food for yourself — the system will reserve it.
  - **Limit** = You can eat at most 2 plates — if you want more, you're not allowed.
- 

### CPU & Memory Units

Resource	Unit	Example
CPU	millicores	500m = 0.5 core
Memory	Mi or Gi	128Mi, 1Gi

---

#### Mi = Mebibyte

- **Full form: Mebibyte**
  - **1 MiB = 1,048,576 bytes =  $2^{20}$  bytes**
  - It's close to 1 MB (Megabyte), but not exactly the same.
- 

#### Gi = Gibibyte

- **Full form: Gibibyte**
- **1 GiB = 1,073,741,824 bytes =  $2^{30}$  bytes**
- It's close to 1 GB (Gigabyte), but slightly more.

---

## Why use Mi and Gi instead of MB and GB?

Because:

- **MB (Megabyte) = 1,000,000 bytes** (decimal - base 10)
- **GB (Gigabyte) = 1,000,000,000 bytes**

So, to avoid confusion between binary and decimal sizes, standards like **MiB and GiB** were introduced.

## Example on Kind Cluster

Let's define a Pod with CPU and memory requests/limits.

---

### 1 Create a Kind cluster

```
bash
-----
kind create cluster --name res-demo
```

---

### 2 Create a Pod with resource limits

```
pod-resources.yaml

yaml
-----
apiVersion: v1
kind: Pod
metadata:
  name: resource-demo
spec:
  containers:
  - name: busybox
    image: busybox
    command: ["sh", "-c", "while true; do echo running; sleep 5; done"]
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"
      limits:
        memory: "128Mi"
        cpu: "500m"
```

---

### 3 Apply the Pod

```
bash
-----
kubectl apply -f pod-resources.yaml
```

Check the pod:

```
bash
-----
kubectl get pod resource-demo
kubectl describe pod resource-demo
```

You'll see the memory and CPU requests/limits under **Containers > Resources**.

---

#### **4 Verify via Metrics (Optional - Requires Metrics Server)**

If you want to monitor actual usage, install the Kubernetes metrics server:

```
bash

kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
```

```
-----
kubectl top pod resource-demo
```

*(Note: metrics-server doesn't come pre-installed on Kind, extra setup needed.)*

---

### **What happens if usage exceeds limits?**

- **If memory exceeds limit** → The container is **killed** (OOMKilled).
  - **If CPU exceeds limit** → The container is **throttled**, not killed.
- 

### **Bonus: Why Requests & Limits Matter?**

- Kubernetes uses **requests** to decide **where to schedule** the pod.
  - **Limits** help prevent one container from hogging all the resources in the node.
- 

### **Cleanup**

```
bash
-----
kind delete cluster --name res-demo
```

simulate an OOM (Out of Memory) error

---

### **Goal**

- We will create a Pod with:
  - **Memory limit = 100Mi**

- Then we'll run a script inside the container that tries to **allocate 200Mi of memory**.
- This will **crash** the container and you'll see OOMKilled in pod status.

---

## Step-by-Step: Simulate OOMKilled in Kind

---

### 1 Start your Kind cluster (if not already running)

```
bash
-----
kind create cluster --name oom-demo
```

---

### 2 Create a pod with low memory limit

```
yaml
-----
# pod-oom.yaml
apiVersion: v1
kind: Pod
metadata:
  name: oom-test
spec:
  containers:
    - name: memory-eater
      image: busybox
      command: ["sh", "-c", "x=$(head -c 200M </dev/zero | tr '\0' 'x'); sleep
300"]
      resources:
        limits:
          memory: "100Mi"
        requests:
          memory: "50Mi"
```

Save as pod-oom.yaml, then apply:

```
bash
-----
kubectl apply -f pod-oom.yaml
```

---

### 3 Watch the pod status

```
bash
-----
kubectl get pod oom-test -w
```

After a few seconds, the pod will restart with STATUS: CrashLoopBackOff.

---

### 4 Confirm it was OOMKilled

```
bash
-----
```

```
kubectl describe pod oom-test
```

Look for something like:

```
makefile
-----
State: Terminated
Reason: OOMKilled
```

---

## Cleanup

```
bash
-----
kubectl delete pod oom-test
kind delete cluster --name oom-demo
```

---

## 2 What happens if a container exceeds its CPU limit in Kubernetes?

- A. It gets more CPU resources automatically
- B. It is throttled and restricted from using more CPU
- C. It crashes immediately
- D. The Pod is terminated

**Answer:** B. It is throttled and restricted from using more CPU

## Which of the following statements is true about memory limits in Kubernetes?

- A. Exceeding memory limit causes throttling
- B. Exceeding memory limit has no effect
- C. Exceeding memory limit causes the container to be killed
- D. Memory limits are optional and cannot be enforced

**Answer:** C. Exceeding memory limit causes the container to be killed

## What is the default behavior if you don't specify resource requests and limits for a container?

- A. The container is not allowed to run
- B. Kubernetes assigns zero resources
- C. The container can use any amount of resources on the node
- D. Kubernetes assigns random values

**Answer:** C. The container can use any amount of resources on the node