



What is a rolling update in Kubernetes?

A **rolling update** in **Kubernetes** is a deployment strategy used to **gradually update** your application with **zero downtime** (or minimal downtime), by **replacing pods one at a time** with the new version.

How it works:

When you update a Deployment (for example, changing the container image version), Kubernetes:

1. **Creates a new pod** with the updated configuration.
2. **Waits** until the new pod is running and healthy.
3. **Terminates one of the old pods.**
4. **Repeats** the process until all old pods are replaced.

This way, your app **remains available** throughout the update.

Benefits:

- **No downtime** during updates.
- **Easy rollback** if something goes wrong.
- Supports controlled rollout using settings like:
 - `maxUnavailable`: max number of pods that can be down during update.
 - `maxSurge`: max number of extra pods that can be created during update.

Step 1: Create an initial deployment

Save this YAML as `deployment-v1.yaml`:

```
yaml
-----
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
        - name: myapp-container
          image: nginx:1.19 # Initial version
          ports:
            - containerPort: 80
```

Apply it:

```
bash
-----
kubectl apply -f deployment-v1.yaml
```

Check pods:

```
bash
-----
kubectl get pods -l app=myapp

kubectl get pods -l app=myapp -o wide
```

Check version

```
kubectl describe deployment myapp-deployment | grep Image
```

Step 2: Trigger a rolling update

Now update the image to a newer version (nginx:1.21).

You can do this using the `kubectl set image` command:

```
bash
-----
kubectl set image deployment/myapp-deployment myapp-container=nginx:1.21
```

Step 3: Watch the rolling update

To see the update in progress:

```
bash
-----
```

```
kubectl rollout status deployment/myapp-deployment
```

This will show you how Kubernetes replaces the old pods one-by-one with new ones running the updated image.

Check the pods again:

```
bash
-----
kubectl get pods -l app=myapp -o wide
```

You'll see the old version getting terminated and new ones taking over.



Step 4: Rollback (if needed)

If something goes wrong, you can easily rollback:

```
bash
-----
kubectl rollout undo deployment/myapp-deployment
```

1. What is the main purpose of a rolling update in Kubernetes?

- A) To delete all old pods at once and start fresh
- B) To update the cluster nodes automatically
- C) To update pods gradually with minimal or no downtime
- D) To restart the Kubernetes API server



Correct Answer: C

2. Which Kubernetes object supports rolling updates by default?

- A) Pod
- B) Service
- C) Deployment
- D) Namespace



Correct Answer: C

3. During a rolling update, what does Kubernetes do first?

- A) Deletes all old pods
- B) Creates new pods with the updated version
- C) Updates the kubelet
- D) Scales down the deployment to zero

✓ Correct Answer: B

4. What does the **maxSurge** parameter define in a rolling update strategy?

- A) The number of pods that can be updated in parallel
- B) The maximum number of unavailable pods during the update
- C) The number of extra pods allowed during the update
- D) The delay between pod updates

✓ Correct Answer: C