

# Explain Taints & Tolerations with a Simple Analogy and also provide real use cases

## **Taints & Tolerations - Simple Analogy**

## **School Classroom Analogy:**

- Imagine a **classroom** (Node) where students (Pods) can come and learn.
- But the classroom has a **signboard (taint)** that says:
  - "No students allowed unless they are vaccinated."
- Most students are **not vaccinated**  $\rightarrow$  **X** they are **not allowed** in.
- But some students carry a vaccination card (toleration)  $\rightarrow \bigvee$  they can enter.

### So:

- Taint = Node says "don't allow anyone unless they match this rule."
- Toleration = Pod says "I'm okay with that rule, let me in."

## In Simple Words:

Concept Meaning (Plain English)

**Taint** Node puts up a "DO NOT ENTER unless you meet this condition" sign

Toleration Pod shows a "PASS" that says "I meet the condition, let me in"

## Summary

**Term** Where is it applied? What does it do? **Taint** On Node Blocks Pods unless tolerated **Toleration** On Pod Allows Pod to go to tainted Node

## Real-Life Use Cases of Taint & Toleration

- 1. Dedicated Nodes for Critical Apps
  - You taint a node: env=critical:NoSchedule
  - Only critical apps (Pods with matching toleration) run there.
- 2. GPU Nodes
  - You taint GPU nodes to ensure only GPU-requiring workloads get scheduled.
- 3. Node Maintenance
  - Temporarily taint a node to stop new Pods from being scheduled during upgrades.

## Demo



- Docker is installed
- kubectl and kind are installed and set up



# Step-by-Step Demo: Taint + Toleration on Kind Cluster



## Step 1: Create a Kind Cluster

```
bash
kind create cluster --name taint-demo
```

## Verify the cluster:

bash kubectl get nodes

## Let's say it shows a node:

pgsql

NAME STATUS ROLES AGE VERSION

## Taint the Node

bash

kubectl taint nodes taint-demo-control-plane key=value:NoSchedule

**☑** Now the node says: "Don't schedule pods unless they tolerate this taint:

key=value:NoSchedule"

-----

## **Step 2: Create a Pod WITHOUT toleration**

```
yaml
-----
# pod-without-toleration.yaml
apiVersion: v1
kind: Pod
metadata:
   name: no-toleration-pod
spec:
   containers:
   - name: nginx
    image: nginx

   image: nginx

kubectl apply -f pod-without-toleration.yaml
kubectl describe pod no-toleration-pod
```

Pod will stay in Pending state — can't be scheduled.

## **Step 3: Create a Pod WITH Toleration**

```
yaml
-----
apiVersion: v1
kind: Pod
metadata:
  name: tolerated-pod
spec:
  tolerations:
  - key: "key"
    operator: "Equal"
    value: "value"
    effect: "NoSchedule"
containers:
  - name: nginx
    image: nginx
```

kubectl apply -f pod-with-toleration.yaml kubectl describe pod tolerated-pod

This pod will be scheduled successfully!

## Step 4: Clean Up

bash

kubectl delete pod no-toleration-pod tolerated-pod kubectl taint nodes minikube key=value:NoSchedule-

----e end removes the taint.



## Basic Syntax Recap

bash

kubectl taint nodes <node-name> <key>=<value>:<effect>

- -----value> define the *taint condition*
- <effect> decides how strictly this condition is applied:
  - NoSchedule: Do not schedule pods that don't tolerate it
  - PreferNoSchedule: Try to avoid scheduling pods here
  - NoExecute: Evict existing pods too if they don't tolerate it



# **Useful Examples**

## 1. ■ Reserve node for critical apps only

kubectl taint nodes node1 env=critical:NoSchedule

• Node node1 will reject all pods except those with toleration for env=critical.

## 2. Mark node for GPU workloads

bash

kubectl taint nodes node2 hardware=gpu:NoSchedule

• Only pods that need GPO and have the right toleration will be scheduled on hode2.
restriction (PreferNoSchedule)
bash
kubectl taint nodes node3 team=data:PreferNoSchedule
• Kubernetes will <b>try not to</b> schedule pods on node3 unless necessary, unless they tolerate it.
enance Mode (Evict running pods)
bash 
kubectl taint nodes node4 maintenance=true:NoExecute
• All running pods will be evicted unless they have toleration for maintenance=true.
m Key-Value for isolating a team
kubectl taint nodes node5 team=frontend:NoSchedule
• Only Pods with toleration for team=frontend can go to node5.
Taint
If you want to remove the taint, just add a – at the end:
bash
kubectl taint nodes nodel env=critical:NoSchedule-

## What is the purpose of taints in Kubernetes?

- A) To increase pod resource limits
- B) To prevent all pods from running on a node
- C) To repel pods from a node unless they tolerate the taint
- D) To automatically restart pods when a node fails

$\checkmark$	Answer:	$\mathbf{C}$

