




## What critical data should be included in a Kubernetes backup strategy?


-  What Needs to Be Backed Up?
    1. **Etcd Database** (the brain of the cluster)
      - Stores **all cluster state**: deployments, secrets, configmaps, etc.
    2. **Persistent Volumes (PVs)** if your apps store data
    3. **Kubernetes YAML Manifests** (your application specs)
- 

### Backup Methods

#### 1. Backup etcd (Control Plane only)

If you have access to the control plane (self-managed clusters):

```
bash
-----
ETCDCTL_API=3 etcdctl \
  --endpoints=https://127.0.0.1:2379 \
  --cacert=/etc/kubernetes/pki/etcd/ca.crt \
  --cert=/etc/kubernetes/pki/etcd/server.crt \
  --key=/etc/kubernetes/pki/etcd/server.key \
  snapshot save /backup/etcd-snapshot.db
```

-  This creates a snapshot of the cluster state.

#### What is **etcdctl**?

- **etcdctl** is the **command-line tool** for interacting with an **etcd** database.
  - Think of it like the **kubectl** for Kubernetes, but this one is for etcd specifically.
- 

#### What is **ETCDCTL\_API=3**?

- It's an **environment variable** that tells the **etcdctl** command to use **API version 3**.
- etcd has two main APIs (v2 and v3), and v3 is the current and recommended version.

- Without setting this, some commands may fail or behave unexpectedly.

---

## 2. Use Velero (for cluster + volume backups)

Velero is a powerful, open-source backup tool.

```
bash
-----
# Install Velero (example with AWS backend)
velero install \
  --provider aws \
  --bucket my-backup-bucket \
  --plugins velero/velero-plugin-for-aws:v1.5.0 \
  --backup-location-config region=us-east-1

# Backup all resources
velero backup create full-backup --include-namespaces '*'

# View backups
velero backup get
```

You can even schedule recurring backups!

---

## 3. Export YAMLs (basic method)

For simple or small clusters, export manifests:

```
bash
-----
kubectl get all --all-namespaces -o yaml > all-resources.yaml
kubectl get configmap,secrets --all-namespaces -o yaml > configs.yaml
```

Later, you can restore:

```
bash
-----
kubectl apply -f all-resources.yaml
```



## Restore Methods

### 1. Restore etcd snapshot (self-managed)

```
bash
-----
ETCDCTL_API=3 etcdctl snapshot restore /backup/etcd-snapshot.db \
  --data-dir=/var/lib/etcd-from-backup
```

Update your Kubernetes config to point to the new etcd data dir.

---

## 2. Velero Restore

```
bash
```

```
-----  
velero restore create --from-backup full-backup
```

---

### 3. Re-apply YAML files

```
bash  
-----  
kubectl apply -f all-resources.yaml
```

This brings back your deployments and services.

---



### Tips

- **Automate backups** (cron jobs, Velero schedules)
  - **Test restores** regularly (especially on staging)
  - Use **external storage** (S3, GCS) for backup reliability
  - Back up **secrets and configmaps** too!
- 

### Amazon EKS (Elastic Kubernetes Service):

---



### How to Backup and Restore EKS Cluster

#### Backup:

- Use **Velero** to back up Kubernetes resources and EBS volumes to **Amazon S3**.
- Optionally, use **AWS Backup** for EBS snapshots.

#### Restore:

- Use **Velero restore** to recreate resources from S3.
- Make sure IAM roles and S3 permissions are properly set up.

Which of the following is the best combination to fully recover a Kubernetes cluster after a disaster?

- A) Pod logs, kubelet binaries, ingress rules
- B) Etcd backup, Persistent Volumes, and YAML manifests
- C) Node IPs, load balancer settings, and Helm charts
- D) Docker images, CRDs, and container logs

**Correct Answer:** B) Etcd backup, Persistent Volumes, and YAML manifests

If the etcd database is lost and no backup is available, what is the most likely consequence?

- A) Nodes will reboot automatically
- B) Only network policies are lost
- C) The cluster will lose all its configuration and state
- D) Persistent Volumes will be unaffected

**Correct Answer:** C) The cluster will lose all its configuration and state